



In Validations We Trust? The Impact of Imperfect Human Annotations as a Gold Standard on the Quality of Validation of Automated Content Analysis

Journal:	<i>Political Communication</i>
Manuscript ID	Draft
Manuscript Type:	Original Article
Keywords:	Automated text analysis, reliability, validation, Monte-Carlo simulations

SCHOLARONE™
Manuscripts

Abstract

Automated text analysis has become increasingly popular in extant political communication research for analyzing texts. While researchers routinely and conveniently resort to some forms of human coding for validating the results from automated procedures, the actual “quality assurance” of such a “gold standard” often goes unchecked in practice. Against this backdrop, this study first presents a systematic review of articles published in major communication and general social science journals in the past 20 years. The results show that the current practices of validation utilizing manual annotations are far from being acknowledged in the literature, and that the reporting and interpretation of validation procedures differ greatly. In a second step, we assess the connection between the quality of human judgment in manual annotations and relative performance evaluations of automated procedures (against true standards) by relying on large-scale, systematic Monte-Carlo simulations. Results from the simulations confirm the expectation that there is a greater risk of a researcher to reach an incorrect conclusion regarding the performances of automated procedures when the quality of manual annotations used in validation is not properly ensured. Our contribution should therefore be read as a call for a systematic application of high-quality manual validation materials in any social science publication drawing on automated text analysis procedures.

Keywords: Automated text analysis, reliability, validation, Monte-Carlo simulations

In Validations We Trust? The Impact of Imperfect Human Annotations as a Gold Standard on the Quality of Validation of Automated Content Analysis

Automated text analysis methods have become increasingly popular for analyzing texts in the social sciences and in the field of political communication in particular, ranging from analyses of decades of newspaper coverage and parliamentary debates to millions of social media posts. Taking advantage of the fact that ever-growing quantities of text are readily available, while resources to analyze these are usually limited, research in the social sciences nowadays eagerly turns to automated approaches. However, with the growing popularity of such “text-as-data” approaches, the issue of the validity of the results and hence the conclusions drawn from them becomes crucial. Blindly applying automated approaches without proper and careful validation may result in misleading or even plainly wrong findings: a principle famously illustrated by the phrase “garbage in, garbage out.”

To arrive at valid results, text-as-data approaches depend on a proper triangulation of the applied techniques against some gold standard (Grimmer & Stewart, 2013).¹ Typically this is done by using human inputs (“human coding” or “manual annotations”) as a benchmark. This is based on the assumption that humans’ understanding of texts (still) outperforms that of machines and that, *if trained correctly*, humans will make the most correct and valid classifications of texts. Human coded data is thus treated as a gold standard against which the performance of the computer is judged. Yet, “the quantities we seek to estimate from text [...] are fundamentally unobservable” (Lowe & Benoit, 2013, p. 299), and as documented in traditional content-analytic applications (Ennser-Jedenastik & Meyer, 2018; Hayes & Krippendorff, 2007; Krippendorff, 2004), human judgments are in fact no exception to this general rule. In this regard, ensuring an acceptable quality of manual coding

¹ Here, we use the term “gold standard” and “ground truth” largely interchangeably, denoting some forms of “objective” (or intersubjectively valid) data that serve as the reference.

(e.g., proper coder training and achieving an adequate level of intercoder reliability) in traditional manual content analysis is regarded as *the* standard practice in the field. However, as we will argue below, there has been a relative lack of parallel attention to ensuring an acceptable quality in human coding when such manual annotations are later utilized as a gold standard for the validation of automated procedures. When it comes to actual research practices, often the human-coded materials appear to be misused by not being properly evaluated and checked before being utilized as validation materials. Such practice, we argue, would greatly percolate additional errors unbeknown to researchers. Yet, the implications of using such imperfect human judgments as a benchmark in validating the results of automated procedures are, until today, insufficiently addressed. Accordingly, this study addresses both the actual practices of the usage and reporting of gold standards in automated procedures and the possible implications of different qualities of such material. Doing so it provides concrete benchmarks as to the desirable quality of “gold standards” produced by human annotations.

Against this background, in the empirical section we first present a systematic review of articles published in major communication and general social science journals over the past 20 years, showing that the actual practices of systematic validation are far from what is being acknowledged as a standard in the literature. Here we also show that the reporting and interpretation of validation procedures differ greatly across studies, and most importantly, that human-coded materials are frequently misused due to a lack of “quality” assessment before being utilized in validations. More importantly, in a second step we argue that using such imperfect human coding as the benchmark has systematic consequences for evaluations of the performances of the different automatic procedures. We assess this previously unexplored connection by relying on a large-scale Monte-Carlo simulation. To our knowledge, this contribution is among the very first to provide thorough, systematic evidence pertaining to a well-known, yet extremely scarcely discussed topic. Our contribution should

be read as a call for a systematic application of high-quality human coding for validation procedures in automated content analysis. To this end, it provides benchmarks for the combination of different levels of the quality of gold standard materials, and warns against improper use of human coding as the benchmark in demonstrating the performances of an automated text analysis approach.

The Use of Human Annotations in Validation of Automated Content Analysis

We define “automated content analysis” (or automated text analysis) as the collection of content-analytic approaches that utilize automated methods of coding a large amount of textual data, in a way that the coding itself (i.e., the text classification) is not performed manually but conducted through predefined computational algorithms (Grimmer & Stewart, 2013), following the standard technique often employed in political communication research.² While the usage of the term “automated content analysis” in general encompasses a wide variety of forms (e.g., Grimmer & Stewart, 2013; Hopkins & King, 2010; Krippendorff, 2013), here we focus two broad forms of automated content analysis -- a *dictionary approach* and *supervised machine learning* methods. We choose to do so since many of the current applications in political communication heavily rely on those two broad approaches (e.g., Aldering & Vliegenthart, 2016; Boomgaarden & Vliegenthart, 2009; Burscher et al., 2015, Burscher et al., 2014; González-Bailón & Paltoglou, 2015; Rooduijn & Pauwels, 2011; Scharkow, 2013; Young & Soroka, 2012).

For automated content analysis methods, the discussion of “validation” has traditionally evolved around the logic of establishing correlative or convergent validity

² Whereas dictionary-based or supervised learning methods assume already well-defined categories, unsupervised methods such as LDA topic modeling generally aim to inductively “discover” new classifications without any human input. Since our primary interest lies in the interplay of human input and machine output, we do not consider unsupervised methods here. Also, our definition (inevitably) excludes automatic approaches of merely *acquiring* data, or traditional manual content analysis but make use of automated procedures in tasks other than the actual coding process (such as in data entry or data management: e.g., Lewis, Zamith, & Hermida, 2013).

(Krippendorff, 2008), mainly by comparing a given measurement with another measurement of the same concept. Therefore, applications of validation procedures in automated content analysis have typically relied on some sort of human input measuring the same concept serving as a benchmark, or a gold standard, as discussed above. For instance, in dictionary-based approaches, the actual coding process itself – using the existing dictionary – does not involve any human inputs. Instead, the use of manual coding in dictionary approaches involves *post-hoc* comparisons of the derived results against manual coding. A recent work by Rooduijn and Pauwels (2011) on the automated measurement of populism in election manifestos well exemplifies this strategy. They show that dictionary-based automatic coding of “populism” categories in party manifestos produces essentially very similar results compared to manual coding of a trained researcher. Similarly, Young and Soroka (2012) compared the manually coded newspaper content against results based on the Lexicoder Sentiment Dictionary (LSD), and found that results are largely comparable to each other.³ While such post-hoc validations could ensure the soundness of conclusions drawn from dictionary approaches, it is generally rare to see a validation of results *after* such classification tasks (yet for notable exceptions, see Muddiman, McGregor, & Stroud, 2018; Young & Soroka, 2012).

In supervised methods, the role of human input is more central in fine-tuning the algorithm (i.e., an “inner” cross-validation). As exemplified in Scharkow’s work (2013), a researcher typically produces manually annotated sample materials, or a training set, and the algorithm is then trained on sample material in order to develop statistical models (that

³ However, a great deal of labor-intensive human input is usually required when building and constructing a well-defined dictionary (Muddiman, McGregor, & Stroud, 2018; Young & Soroka, 2012). Due to its labor-intensive nature, recent applications in this area increasingly turn to “crowdcoding”, where the manual labor of highly trained yet few numbers of human coders are replaced with a large number of (little or untrained) crowd-coding workers (Haselmayer & Jenny, 2017; Lind, Gruber, & Boomgaarden, 2017).

effectively aim to reproduce implicit coding rules of human coders) to predict and classify unseen materials (i.e., the test set). For instance, Burscher et al. (2014) have developed a supervised machine-learning algorithm to automatically classify generic media frames in news articles, based on a random subset of data that were manually coded by trained coders.⁴ Yet due to inherent resource constraints, such validations typically rely on only a very small subset of held-out samples that have to be coded *both* by human coders and by the trained algorithm. As such, post-hoc validation as described in dictionary-based approaches (i.e., an “outer” cross-validation) appears to be also rarely offered in practice.

The Myth of Perfect Human Coding in Validation of Automated Content Analysis

Regardless of its specific orientations, the use of human coding as a gold standard is often regarded as *the* principal method of ensuring the validity and soundness of conclusions derived from the automated procedures (DiMaggio, 2015; Grimmer & Stewart, 2013). The purpose of such validation against a gold standard is, as Krippendorff (2008) notes, “to confer validity on the otherwise uncertain research results” (p. 6). This logic requires the chosen benchmark (i.e., manual annotations by human coders) to be of objective and unquestionable truth – *or at its minimum*, to be sufficiently intersubjective to establish a firm common ground (Krippendorff, 2008). Yet, as much of the traditional manual content-analytic methodology literature suggests (e.g., Ennser-Jedenastik & Meyer, 2018; Hayes & Krippendorff, 2007; Krippendorff, 2004), manual coding often produces unreliable judgments under a lack of proper coding instructions or coder training, and especially so when the judgment in hand requires a nontrivial degree of inferences and subjectivity to classify latent information (Krippendorff, 2013). For this reason, there has been relatively

⁴ In unsupervised methods, the use of human-coding as a benchmark is a common practice as well. For instance, Lowe and Benoit (2013) directly compare direct human-coding based party position scaling (which serves as a benchmark) against unsupervised scaling methods; they found that the proposed scaling methods can produce largely similar results on par with human judgments.

little disagreement regarding the need for proper “quality assurance” in the form of developing unambiguous coding categories and better coder training to achieve an acceptable level of intercoder reliability (Krippendorff, 2004, 2013).

Yet when researchers apply automated procedures, the “quality” of manual annotation often goes unchecked even if researchers routinely and conveniently resort to some forms of human annotation in validating their results (if at all). Instead, researchers are often put too much trust in such naïvely-coded human annotations in typical validation procedures. Given the importance of proper validation in text-as-data approaches (Boumans & Trilling, 2016; Grimmer & Stewart, 2013), a seemingly widespread practice of conveniently utilizing quick and easy manual annotation without proper quality assurance in typical automated text analysis applications is especially surprising. “[W]henever ... principles by which humans generate ratings are heterogeneous across raters” (DiMaggio, 2015, p. 4), any automated procedures that are evaluated upon such imperfect manual annotations would be systematically biased to the degree that can be found in such imperfect human judgments.

The use of (potentially) imperfect, low-quality manual annotation as a benchmark for automatic techniques may have at least two systematic consequences. First, while automated methods themselves are perfectly reliable, imperfect human judgments on validation materials essentially makes the ultimate “target” of such reliable measurements radically deviate from the true target of inference, making them “reliably wrong” on-target. Second, and relatedly, systematically flawed human judgment can introduce unknown errors to algorithms, leading to biased conclusions against a true standard. Therefore, using suboptimal manual annotation on validation materials makes it harder to evaluate the relative trustworthiness of such validation procedures. However, most of the empirical studies appear to pay insufficient attention to this issue. Indeed, we know only one existing study that suggests some tentative evidence of the relationship between quality of human coding (on

validation materials) and machine-based classification accuracy (Burscher et al., 2014).⁵

In sum, there are good reasons to suspect a systematic relationship between the use of low-quality manual annotation and the resulting bias⁶ and errors regarding the ultimate conclusions drawn from automated content analysis. Many prior contributions on this topic – both theoretically and empirically – stress a need for a proper validation of applied techniques (e.g., Grimmer & Stewart, 2013; González-Bailón & Paltoglou, 2015; Hopkins & King, 2010) and generally speaking such issues are gaining more attention in various applied research domains (e.g., such as in corpus annotations: Hovy & Lavis, 2010; Lease, 2011). However, we do not know much about where the field in general stands in terms of standard validation practices in actual research and the use of imperfect human coding in particular.

Against this backdrop, we first present a systematic review of relevant literature. While the first empirical section should provide an overall picture of how validation is typically approached in automated content analysis in social science practice, at the same time it serves to benchmark our subsequent simulation study.

Systematic Review of the Relevant Literature

Sample and Procedures

We have identified relevant studies using the EBSCOhost databases “Communication & Mass Media Complete,” “Humanities Source,” and “SocINDEX with Full Index,”⁷ querying all titles, abstracts, and keywords using the following Boolean search string:

⁵ In a recent study, González-Bailón and Paltoglou (2015) compare five available sentiment dictionaries against human annotations, yet they do not directly deal with the implications of imperfect reliability in human coding. Scharnow and Bachl (2017), the only one of existing studies that examines the consequences of imperfect reliability in human coding, mainly deal with its implication on “linkage analysis,” but not on automated content analysis.

⁶ This is a different issue from much-discussed “representational bias” (especially within the context of an algorithmic bias), but rather is more closely related to “quality control” (QC) issue in natural language processing literature (Lease, 2011). We use the term “bias” here to refer to a degree of discrepancy between true classification membership that could have been obtained by (theoretically) perfect data and that from low-quality data.

⁷ This was to ensure that we mainly focus on political “communication” applications.

("computer assisted" OR "automated" OR "automatic" OR "computational" OR "machine learning") AND ("content analysis" OR "text analysis"). This resulted in a total of 192 English-language journal articles published between January 1, 1998 and November 7, 2018. Among them, 119 articles were determined as not relevant (e.g., non-empirical overviews/introduction articles, qualitative analyses, studies using unsupervised methods, or simple keyword frequencies, etc.)⁸ and 7 articles were either duplicates or could not be obtained as full texts. These articles were excluded from further analyses. Using the remaining 73 articles, we systematically examined whether extant applied studies a) employed any empirical validation of their primary findings, b) if so, whether they used human coded materials as a benchmark, and c) if so, whether intercoder reliability and other methodological details were adequately reported.

A total of five highly trained coders tested the initial coding scheme by independently coding 10 sample articles (approximately 5% of the total retrieved sample) and collectively discussed any coding problems and disagreements. Coding instructions were iteratively revised until the coding schemes would produce reliable results. Intercoder reliability (based on Krippendorff's alpha) above 0.75 was ensured for each of the variables coded (see the online appendix for detailed information regarding coded variables, including coding instructions and detailed reliability estimates).

Results

The results of the systematic review of the literature are presented in Table 1 below. Among 73 articles, a total of 55 studies used dictionary approaches while 18 used supervised machine learning methods.⁹ From the papers using a dictionary approach, only about half of

⁸ Here, we do not consider a simple keyword-frequency based study (e.g., simply counting the number of occurrences of a keyword in a given text, but not actually classifying the documents based on such frequency) since human inputs play no role other than compiling the keyword list itself.

⁹ Among excluded studies, only 15 studies have used unsupervised learning or other forms of automated content analysis, suggesting dictionary-based or supervised machine learning applications

IN VALIDATIONS WE TRUST

1

1
2
3 them referred to some sort of manual gold standard in their text, and typically relied on 3.4
4
5 coders with on average a total of 532.86 annotation materials, yielding a little over than 150
6
7 manual annotations per coder. Among them, only 16 percent provided any measures of inter-
8
9 coder reliability for such materials: Two studies relied on Krippendorff's alpha whereas three
10
11 studies reported either percentage agreement or Holsti agreement measure (in the remaining
12
13 studies, we found other measures such as Scott's Pi or correlation coefficients). Notably,
14
15 shares are much higher when it comes to articles using supervised machine learning. Here, a
16
17 total of 16 studies (out of 18 studies, or 88.8%) referred to manual gold standard materials,
18
19 and among them, 10 studies typically relied on 5.7 coders. Excluding one outlier (a study
20
21 based on more than 20,000 manually coded materials), they on average used a total of 1,640
22
23 annotation materials, yielding approximately 300 manual annotations per coder. Six studies
24
25 did not report the number of coders and/or the total size of the validation dataset, making it
26
27 impossible to judge the soundness of the applied validation procedures. In sum, the most
28
29 common measures of inter-coder reliability (if reported) were percentage agreement (38%, N
30
31 = 7) and Krippendorff's alpha (33%, N = 6). Among 73 articles we examined, only 9 studies
32
33 (which is about 20% of studies that referred to some human-coded gold standards) used
34
35 Krippendorff's alpha, with an average value of 0.68.
36
37
38
39
40
41

42
43 Regarding the validation of automated approaches, results were very similar. Only
44
45 around 40 percent of the papers using a dictionary approach also reported some validation
46
47 measures, compared to 67 percent of the papers relying on supervised machine learning. The
48
49 most commonly used measures of validity were indeed the widely accepted measures of
50
51 precision (13 cases, M = 0.74) and recall (9 cases, M = 0.60), yielding an average F1 score of
52
53 0.63 (10 cases, based on either directly reported F1 scores or indirectly calculated F1 scores
54
55 from reported precision and recall). However, other metrics were widely used as well in
56
57
58

59
60 are much more frequently used in social sciences in general.

reporting validation. This is somewhat disconcerting insofar as these tend to be indeed either the intercoder reliability measures (e.g., Holsti, Cohen’s Kappa, Krippendorff’s alpha) or correlation coefficients, both of which are not designed for the validation of a given automated procedure.¹⁰ Interestingly, there were only three studies which reported *both* Krippendorff’s alpha in human-coded materials (which signals the proper coder training and sufficient quality assurance of manual annotations) and proper validation metrics.

Generally, the results of this review reveal that reported measures of validation and the quality of human-coded data used are far from what is being acknowledged as best-practice in the extant literature. While we believe the proper use of manual annotations in automated approaches should comprise both sufficient quality assurance of and the proper use of validation metrics based on such manual annotations, the fact that only a total of three studies (about 4% of all studies) satisfy such criteria reveals a striking lack of attention to this issue when it comes to actual research practice.

-- Table 1 About Here --

A Monte-Carlo Simulation Study

The findings provided above reveal that there is still strikingly little consistency in *whether*, and if so, *how* validation is approached and reported. About half of the studies did not report *any* validation metrics when relying on automated methods. Even when they did, metrics related to (human-coding based) validations and their qualities were generally not consistently reported, and often limited in providing actual methodological details. In order to systematically evaluate the implications of such improper use of human-coded materials in

¹⁰ Generally speaking, intercoder reliability assesses the extent to which two or more coders classify a given content into a “same” category. Such a classification inevitably collapses the *true positive* and the *true negative* into one category (“same”), while the *false negative* and the *false positive* are being lumped into another category (“different”). At best, it can only approximate the classification “accuracy” (defined as $[TP+TN]/[TP+TN+FP+FN]$) under the assumption that only one of the coders *always* produces true classification (the gold standard). Nevertheless, classification accuracy is not a good measure of classification performance when there is a high class-imbalance in the dataset -- which is often the case in many empirical applications.

IN VALIDATIONS WE TRUST

1

the validation of automated procedures, we set up an extensive set of Monte-Carlo (MC) simulations. MC simulation offers a convenient yet flexible tool for systematically evaluating the relative bias (against the true standard) and coverage of a given statistic under certain scenarios (Leemann & Wasserfallen, 2017; Scharkow & Bachl, 2017). A set of replication codes for this manuscript can be found at [redacted for a review].

We designed our procedures in a way that largely mirrors the typical approaches in this area by breaking down our approach into three stages – data generation, human coding, and automatic classification – where we systematically varied the intercoder reliability of the “gold standard” material, along with a number of related factors such as the number of coders and the number of manual annotations per coder (see below). Then we systematically compared different scenarios in terms of their classification accuracy and F1 scores based on the “true” standard (i.e., a quantity of interest that is typically unknown to researchers) in order to illustrate how different practices of human coding in automated content analyses affect the relative trustworthiness of conclusions drawn from such results.

Data Generation Stage

We create data (e.g., textual data to be analyzed by a researcher) with the “true” outcome value of interest; the goal of any quantitative text analysis method is to somehow approximate this true value of y , either by human coding, machine algorithms, or some combinations of both. For the data generating process, we set the true value of y to be randomly generated from three hypothetical independent variables (x_1 , x_2 , and x_3), the values of which were also randomly sampled either from a multivariate normal distribution (for supervised learning) or from a categorical distribution (for the dictionary approach – see online appendix for details). This ensures that the results of our simulations are not completely deterministic, as well as not analytically driven to arrive at our conclusion.

We assume the size of the text data is sufficiently large to warrant an automated

approach. As such, the data in question (hypothetically) covers 10 news articles per day per each of 10 outlets, spanning a total of 20 years. Accordingly, each single simulation run is set to generate 730,000 observations of media content data (10 x 10 x 365 x 20).

Human Coding Stage

In a typical content analysis, at least two or more trained human coders are assigned to a small set of sample documents, and independently code such documents. This process of coder training is then repeated until the satisfactory level of intercoder reliability is achieved (typically Krippendorff’s alpha equal to or greater than 0.7). Once a coder’s reliable understanding and application of the coding instructions is ensured by the training, validation materials are often divided evenly and annotated only by a single coder (Grimmer, King, & Superti, 2018). This means that any potential remaining coder idiosyncrasies are no longer an issue given the satisfactory level of intercoder reliability. However, when the quality of manual annotations is not sufficiently ensured, such idiosyncrasies may systematically bias the procedures to an unknown degree.

Following this logic, we specified the following factors that may affect the “quality” of the gold standard (i.e., manual annotations by human coders) and therefore evaluations of the overall performance of automated algorithms. Those factors include: the number of human coders ($k = 2, 4, 7, 10$), the levels of intercoder reliability (Krippendorff’s $\alpha = 0.5, 0.6, 0.7, 0.8, 0.9$), as well as the number of annotations per coder ($n = 50, 100, 250, 500$). While in any typical real-world application a researcher has to make practical and logistical decisions on these factors (typically guided by a researcher’s resource constraints), these factors are indeed crucial in terms of properly ensuring the acceptable quality of manual annotations produced by human coders. These numbers were chosen to reflect typical procedures and their common variations, as can be seen in our literature review above.

In all scenarios, human coders classify a given observation as “1” (e.g., a text contains

IN VALIDATIONS WE TRUST

1.

the quantity of interest, such as a certain actor, frame, or tonality) or “0” (e.g., does not contain this quantity of interest). This human coding (\hat{y}) can be, in principle, either correct or incorrect against the (unknown) true value of y , therefore behaviors of human coders were modeled by a Binomial distribution with varying probability of successfully categorizing the true data. The heterogeneity in each of the coders’ coding behaviors (e.g., conditioned by expertise) was also modelled by a beta distribution with varying shape (using predefined hyperparameters for each target reliability level), which effectively enables us to simulate a situation where some coders produce more “correct” judgments ($\hat{y} = y$) whereas other coders produce more “false” judgments ($\hat{y} \neq y$). Yet the overall human annotation patterns at the chosen level of a beta distribution parameters were ensured to produce an associated target reliability coefficient.¹¹ Depending on the specific application (dictionary-based vs. supervised learning), the “hand-coded data” from this stage is then used either for post-hoc validation or as an input for the later automated classification.

Algorithm-based Classification and Validation Stage

In this final stage, a researcher uses a certain algorithm to predict the values of (\hat{y}) in each of the observations. For this purpose, we set up three different classification algorithms – binomial GLM, Naïve Bayes, and a bag-of-words dictionary approach – in simulations.

In a typical real-world scenario, researchers would validate the results obtained from the automated analysis based on a manually coded test set (produced in a human coding stage) in a post-hoc manner (for a dictionary approach), or use such manually coded data to develop prediction algorithms (for supervised machine learning, such as binomial GLM or Naïve Bayes), evaluating precision, recall, and F1 scores against such human-annotated

¹¹ It is important to note that, in reality, untrained manual coders often experience “coding drift” over time (i.e., a low *intra-coder* reliability), and this is often correlated with the low intercoder reliability (which signals a lack of a proper coder training). However, since we are only simulating the variability of intercoder reliability alone, the intra-coder reliability factor was not considered in our simulation. We return to this point later in the discussion section.

IN VALIDATIONS WE TRUST

1

materials (hereafter “observed” classification performance). Yet more importantly, since we are effectively simulating textual data, we can also systematically evaluate the *true* classification performance of automated procedures as well, and then further compare this true performance with the “observed” performances from scenarios using varying qualities of human-annotated gold standard materials. Admittedly, this would be impossible for practical applications since the true value of y would never be known in practice. However, by doing so, we can precisely estimate the relative bias under the imperfect human coding as “gold standard” and how the overall true accuracy is adversely affected based on which factors.

The final Monte Carlo simulation used 4 (number of human coders, k) \times 5 (target Krippendorff’s alpha levels in human annotations in producing validation data) \times 4 (N of annotation per each human coder in producing validation data) \times 3 (classification algorithms) full factorial design with 1,000 replications per scenario ($N = 240,000$).

Simulation Results

As an initial model check, we first look at the overall classification accuracy against the true value (true classification performance) as a function of the target reliability in validation materials. This is to check whether our simulation setup can indeed correctly reproduce common patterns of results in extant studies, ensuring the validity of our setup and its results. In Figure A1 of the online appendix, the overall prediction accuracy against the true value (defined as the sum of true positive and true negative cases over all cases) is reported, along with their 95% confidence intervals, for every combination of the experimental factors in the MC simulations. Using the overall accuracy in our simulation as the reference point, the first four and the second four sets of panels in Figure A1 make it clear that the reliability level of the training material has a nontrivial benefit in improving the accuracy of predictions based on automated procedures, especially for ML methods. Indeed, this result is expected since the ML methods take the human input as the basis for developing

IN VALIDATIONS WE TRUST

1

the classification algorithm. Therefore overall accuracy of the final classifier is dependent upon the size of human input and the quality thereof. In contrast, the last four sets of panels in Figure A1 show that the typical Bag-of-words application does not benefit from improved (post-hoc) human coding. Yet again, this is an expected result, since the performance of a given dictionary itself does not depend at all on *post-hoc* validation. Overall, Figure A1 shows that our setup is capable of correctly reproducing a common pattern.

Next, we examine indirect consequences of relying on imperfect, less-than-desired quality manual annotations as a benchmark in evaluating the performances of automated procedures. Typically, researchers rely on a small fraction of held-out human-coded materials at this stage, deciding whether the (observed) overall accuracy or classification performance is good enough to proceed to the next stage. This decision has to be based on some *a priori* chosen threshold value: if validation (based on human coding) is purportedly not satisfying enough to pass such a threshold level, additional steps are sought to improve the quality of automated procedures (e.g., re-training algorithms, or changing the dictionary, etc).¹² The primary interest of such validation lies in extrapolating the observed level of classification performance (based on human-annotated validation materials) to the level of classification performance *that could have been observed* under the perfect, “true” quantity of interest. In other words, the “observed” classification performance based on manual annotations serves as a proxy, or an estimate for the true, yet often unknown, classification performance against the true standard. An interesting question here is thus how well the “observed” classification performance reflects the true classification performance when researchers use imperfect manual annotations, and how large or small a potential bias is compared to the true scores.

In order to illuminate this issue, we divide our simulations into four mutually

¹² Here, we do not consider a scenario where a researcher decides to improve the quality of human coding. This is based on the consideration, as to our argument being advanced here, that a researcher (often incorrectly) assumes that human coding is perfectly reliable and valid.

IN VALIDATIONS WE TRUST

1

exclusive categories as in Table 2 based on the cross-tabulation of “observed” F1 scores (against human-annotated validation materials) versus “true” F1 scores (against the true values of y). We first present the proportion of simulation cases which incorrectly conclude about the overall classification quality based on the “observed” classification performance, using the cutoff value of F1 score = 0.6311 (i.e., the average F1 score reported in the studies reviewed above). Second, we further present the degree of bias, or “error” — defined as the $F_{1_{validation}}/F_{1_{true}}$, where the F1 score is a weighted average of the precision and recall — which captures the degree of under- or over-estimation of true F1 scores against observed F1 scores. In other words, it shows to what degree observed F1 scores using (potentially) imperfect human annotations deviate from the true F1 scores based on “true” standard.

– Table 2 and Figure 1 Here –

Figures 1 presents our simulation results using Naïve Bayes classifier as an illustrative case. As can be seen in Figures 1, it appears that utilizing more “high-quality” manual annotated materials for validation has obvious and discernible consequences in the evaluation of classification quality of the automated procedures. Among 1000 replications of each scenario, all of our experimental factors appear to decrease the decision error rates in using observed level of F1 scores to approximate the true F1 score level. The leftmost upper panel in Figure 1 shows that worst-case scenarios -- using only two coders with a handful of texts (annotations per coder $N = 50$) that are low in reliability (K alpha of 0.5) -- lead to approximately 30% of instances that incorrectly over- or underestimate the true F1 scores. Under the same combination of the number of coders and number of annotations per coder, improving intercoder reliability at best marginally decreases the overall percentage of decision errors. Yet as either the number of coders or the number of independent annotations per coder start to increase, we see that the total proportion of cases that incorrectly estimate the true F1 scores start to decrease substantially (the upper panel of Figure 1). Indeed, when

IN VALIDATIONS WE TRUST

1

the total size of manually annotated validation materials exceeds 1,000 observations, the percentages of decision error is less than 10% (based on all the cases of Type I and Type II error combined), although there also appears to be some variation depending on the specific algorithms being utilized (as appear in the upper panel of Figure A2 and Figure A3 in the online appendix).

The lower panel of Figure 1 displays the point estimate of relative bias (defined as the $F1_{validation}/F1_{true}$, or the degree of under- or overestimation of the true F1 scores based on observed F1 scores), along with their 95% CIs from 1000 replications of each scenario. We see that the distribution of relative bias also starts to converge to unbiased estimates (toward 1, which means observed F1 score is same as true F1 score) as the total number of annotations start to increase, no matter which combinations of total number of coders and number of manual annotations per coder are being used in both, ML scenarios and dictionary approaches. In addition, there initially appears to be no apparent overall main effect of higher reliability of manually-annotated validation materials in terms of reducing the total number of decision errors in ML approaches (as per the upper panel of Figure 1). However, using higher reliability in manual validation materials nevertheless tends to reduce the uncertainties of relative errors, as shown in the lower panel of Figure 1, and especially more so when coupled with a larger size of validation materials. For instance, under a scenario of $N = 100$ validation materials (e.g., 2 coders x 50 annotations per each coder: the first set of bars in lower panel of Figure 1), increasing intercoder reliability from 0.5 to 0.9 reduces the 95% CI range of relative errors from [.95, 1.61] to [.99, 1.53], or a 17.45% reduction in CI range. When the total number of annotations is sufficiently large, such as $N = 5000$ (e.g., 10 coders x 500 annotations per each coder: the last set of bars in lower panel of Figure 1), the same change in reliability from 0.5 to 0.9 reduces the 95% CI range of relative errors from [.92, 1.11] to [.97, 1.03], or 69.89% reduction in CIs. This thus greatly reduces the uncertainties regarding the

IN VALIDATIONS WE TRUST

1

ultimate conclusions one can draw from proposed automated procedures. While we see largely similar patterns for the GLM classifier (Figure A2 in the online appendix), in the dictionary-based applications as presented in Figure A3 in the online appendix, we see a much clearer impact of higher reliability. It appears that observed classification quality (against manually coded validation materials) in these applications tends to underestimate the true classification quality (lower panel of Figure 3). Yet, higher reliability nevertheless greatly reduces the uncertainties regarding the ultimate conclusions based on automated procedures. In sum, this provides a very consistent picture of the impact of “proper” manually coded materials – both in terms of its size and the quality – in validating the conclusions drawn from automated procedures.

Discussion

Based on a thorough discussion of the importance of proper validation in automated text analysis, we attempted to provide a systematic overview of current practices concerning validation procedures of automated content analysis in the social sciences more generally and in the field of political communication in particular. Also, we provide insights into the consequences of using suboptimal-quality manual annotations as a “gold standard” material in automated procedures. The results of the systematic review show that there is strikingly little consistency in *whether* and *how* validation of automated procedures is reported in extant studies. Very often, studies do not report *any* validation metrics when relying on automated methods. Even when they do, validation metrics utilizing human-coding (if any) are generally not consistently reported, and are often limited in providing actual methodological details. They frequently misreport or mis-specify the validation metrics, and often, they do not properly evaluate the quality of human coding when manual annotations are utilized in validation. The results of Monte Carlo simulation procedures also revealed that any practical and logistical factors for standard “quality assurance” – the number of coders, the size of the

training sets produced by each coder, and the “quality cut-off point” (e.g., Krippendorff’s alpha for intercoder reliability) – indeed all produce systematic consequences for the evaluation of the proposed automated procedures. Coupled together, the results from both studies give good reason for concern about the quality (or rather the *validity*) of conclusions drawn from typical automated content analyses. In order to make sensible conclusions, a proper validation of the automated methods is much needed, let alone it needs to adhere to high standards in producing validation materials; yet such proper validation – especially when manual annotations are involved – is rarely done in practice.

To be clear, the current study *does not* make the argument that we should exclusively rely on human validations, or conversely, human validations in general are a problem. Quite the contrary, one of the main points being advanced here is that humans are not perfect, therefore, *proper* validation ensuring the “quality” of manual annotations is essential especially when they are utilized as the gold standard in automated procedures. While the (potentially imperfect) human “gold standard” is often the best we can get, the results of this study suggest that extra caution must be taken. Often, imperfect manual annotations can percolate to machine coding to inferences if a researcher turns a blind eye to such imperfect quality of human judgment. We therefore advise researchers to pay closer attention to the issue of ensuring proper quality of manually coded validation materials.

Some Tentative Recommendations for Best Practices

While there exists no single fool-proof solution applicable to every situation, based on our observations from our systematic review and from our simulation results, here we offer some general recommendations towards the best-practice of utilizing human annotations in validating the automated approaches. First, while we believe not every study with automated content analysis should use human-involved validation, yet *if used*, researchers should adhere to rigorous methodological standards to the degree expected for a traditional manual content

analysis (e.g., Hayes & Krippendorff, 2007; Krippendorff, 2013) in preparing the manually annotated validation dataset. Also, they should fully disclose the methodological details (such as measurement details, coding instructions, coder training, intercoder reliability, etc.) of such validation data, enabling readers to independently judge the soundness of validation procedures, but also for increasing transparency and replicability of the research process.

Second, we recommend researchers to strive to increase the total size of the manually coded validation datasets as large as possible, preferably always more than $N = 1,000$. In our review, we observed that the typical size of the manual annotation dataset ranged from approximately 500 to 600 annotations in total. Yet in our simulation we observed that the risk of potential decision errors (by relying on observed classification accuracy based on low quality validation materials) is substantially high, sometimes up to 30%, with smaller size of the manual annotations. Yet the percentages of decision errors appear to reach the acceptable range -- less than 10% based on all the cases of Type I and Type II error combined -- when the size of the manually annotated validation data is equal or greater than 1000. Here we somewhat arbitrarily suggest the 10% threshold value, yet the combined error rate of 10% roughly represent a balance of considerations between a typical false positive rate (i.e., $\alpha = 0.05$) employed in the field and maintaining a sufficient statistical power (i.e., $1 - \beta$ of 0.95, where β is a false negative rate) as advocated in a recent study (e.g., Holbert et al., 2018). Therefore, we advise researchers to use, *at least*, more than 1000 manually annotated texts as a very minimum criterion. Of course, this number is only a tentative suggestion reflecting a very rough rule of thumb for much simplified cases (i.e., a simple binary judgment involving only one variable), and researchers may make an informed argument of whether they would prioritize which types of errors -- the false positive rate, the false negative rate, or the combined error rates -- hopefully based on evidence presented here. For more complex and subtle judgments (such as ambiguous latent contents or non-binary judgments), it would

IN VALIDATIONS WE TRUST

2

make sense to use much larger size of manual data in validation tasks, although increasing the size of manual annotations in such cases would be a challenging task for its own.

Third, and relatedly, we observed that improving intercoder reliability in human coding *sometimes* offers a large benefit in reducing the magnitude of potential bias in automated classification tasks relative to true, unknown standard. While this seems encouraging for many researchers, at the same time this also warns against the prevalent practice of only prioritizing the intercoder reliability in coder training to reach a minimally “acceptable standard” without considering other factors. Quite contrary, such improvements at best offer marginal gains over potential decision errors especially in cases with very small validation datasets, as evidenced in our simulation study. Recent developments in “crowdcoding” for content analysis (Haselmayer & Jenny, 2017; Lind et al., 2017) could alternatively offer promising ways of scaling up the manual annotation tasks (therefore increase the manual annotations in validation tasks) if resource constraints for using trained coders are high, although the additional issue of quality control for crowdworkers quickly becomes an important issue in this case as well (Lease, 2011; Lind et al., 2017).

Lastly, one should also bear in mind that without proper coder training in improving reliability, human coders often experience substantial “coding drift” over time (i.e., a low *intra-coder* reliability), and this often goes hand in hand with low inter-coder reliability as well. In such a case, the risk of introducing additional random errors due to low intra-coder reliability runs very high. However, since our simulation setup does not take the intra-coder reliability (but only inter-coder reliability) into account, our simulation as a consequence would have produced more “optimistic” results than most of the real-world, low inter-coder reliability scenarios. For that matter, our results should *not* be read as a signal that inter-coder reliability or human coders do not matter at all in validation of automated procedures.

Limitations and Conclusions

Few limitations should be noted. First, the systematic review is limited to prior studies explicitly containing our search terms in their title, abstract, or keyword. Such sample, by definition, does not include all potentially relevant articles. Still, it allows us to make a sensible selection of extant articles that do not only use these methods but also advertise that they do so. In fact, we believe articles that are not part of our sample to be of even lesser rigor in the use and reporting of “gold standards” in validation materials.

Second, in the simulations we only considered a binary classification task and a single dimension of validation metrics — specifically, recall, precision, and resulting F1 scores. While this greatly simplifies our main arguments and (still complex) simulation setups, there are indeed nontrivial numbers of existing applications that go beyond such simple classifications, dealing with numerical forms of predictions (i.e., a document scaling). Yet we reason that our core arguments can be equally applicable to more complex forms of automated content analysis as well. Given the additional complexity and difficulties involved in such non-binary predictions for human coders, achieving acceptable intercoder reliability for manual validation materials in such applications is likely to be even more difficult than in simple binary ones. Hence we suspect the potential problems of “imperfect quality” in human-coded validation materials should be greater, at least identical, in those applications.

Designing a simulation-based study gives us a unique opportunity to observe many potential counterfactual scenarios of the research processes. Such an approach, if carefully designed, allows to robustly explore potentially important variability in research practices and their consequences. Also, one of the core advantages of relying on such a simulation-based approach is that it provides an opportunity of formally checking the sensitivity of one’s findings, or enables one to explore possible counterfactual scenarios. Such clairvoyance, however, comes at a cost — the degree of abstraction and simplification. This simplification is done not only for computational, but also for conceptual reasons. As Scharkow and Bachl

IN VALIDATIONS WE TRUST

2.

(2017) note, “the challenge is to specify a simulation that is simple enough to be comprehensible, yet realistically models the underlying process of interest” (p. 330). In this regard, we have relied on somewhat idiosyncratic and simplified approaches in our simulations. While we surely acknowledge that our setup could have been constructed in a more realistic yet much more complex way, our ability to accurately simulate human behaviors does not necessarily depend on very complex models.

Notwithstanding the aforementioned limitations, we believe that our contribution would further prompt political communication researchers to pay closer attention to the issues of systematic and proper validation of automated content analytic methods. Our contribution should be read as a call for a thorough and systematic application of validation procedures – especially the ones involving manually annotated materials as a “gold standard.” It is worth stressing here that automated content analysis should not be just regarded as a cheap alternative to expensive manual coding, but it also takes – and indeed should take – a good amount of time and resources in designing the study and evaluating its performance. To this end, a little extra effort on designing proper validation does quickly pay off, and eventually would ensure valid inferences and conclusion drawn from such studies.

References

Aaldering, L., & Vliegenthart, R. (2016). Political leaders and the media: Can we measure political leadership images in newspapers using computer-assisted content analysis? *Quality & Quantity*, 50, 1871–1905. doi:10.1007/s11135-015-0242-9

Boomgaarden, H. G., & Vliegenthart, R. (2009). How news content influences anti-immigration attitudes: Germany, 1993–2005. *European Journal of Political Research*, 48, 516–542. doi:10.1111/j.1475-6765.2009.01831.x

Boumans, J. W., & Trilling, D. (2016). Taking stock of the toolkit: An overview of relevant automated content analysis approaches and techniques for digital journalism scholars. *Digital Journalism*, 4, 8–23. doi:10.1080/21670811.2015.1096598

Burscher, B., Odijk, D., Vliegenthart, R., De Rijke, M., & De Vreese, C. H. (2014). Teaching the computer to code frames in news: Comparing two supervised machine learning approaches to frame analysis. *Communication Methods and Measures*, 8, 190–206. doi: 10.1080/19312458.2014.937527

Burscher, B., Vliegenthart, R., & De Vreese, C. H. (2015). Using supervised machine learning to code policy issues: Can classifiers generalize across contexts? *The ANNALS of the American Academy of Political and Social Science*, 659, 122–131. doi: 10.1177/0002716215569441

DiMaggio, P. (2015). Adapting computational text analysis to social science (and vice versa). *Big Data & Society*, 2(2). doi: 10.1177/2053951715602908

DiMaggio, P., Nag, M., & Blei, D. (2013). Exploiting affinities between topic modeling and the sociological perspective on culture: Application to newspaper coverage of us government arts funding. *Poetics*, 41, 570–606. doi: 10.1016/j.poetic.2013.08.004

Enns-Jedenastik, L., & Meyer, T. M. (2018). The impact of party cues on manual coding of political texts. *Political Science Research and Methods*, 6, 625–633.

doi:10.1017/psrm.2017.29

González-Bailón, S., & Paltoglou, G. (2015). Signals of public opinion in online communication: A comparison of methods and data sources. *The ANNALS of the American Academy of Political and Social Science*, 659, 95–107. doi: 10.1177/0002716215569192

Grimmer, J., & Stewart, B. M. (2013). Text as data: The promise and pitfalls of automatic content analysis methods for political texts. *Political Analysis*, 21, 267–297. doi:10.1093/pan/mps028

Grimmer, J., King, G., & Superti, C. (2018). The unreliability of measures of intercoder reliability, and what to do about it. Unpublished manuscript. Retrieved from <http://web.stanford.edu/~jgrimmer/Handbib.pdf>

Haselmayer, M., & Jenny, M. (2017). Sentiment analysis of political communication: Combining a dictionary approach with crowdcoding. *Quality & Quantity*, 51, 2623–2646. doi:10.1007/s11135-016-0412-4

Hayes, A. F., & Krippendorff, K. (2007). Answering the call for a standard reliability measure for coding data. *Communication Methods and Measures*, 1, 77–89. doi:10.1080/19312450709336664

Holbert, R. L., Hardy, B. W., Park, E., Robinson, N. W., Jung, H., Zeng, C., ... & Sweeney, K. (2018). Addressing a statistical power-alpha level blind spot in political-and health-related media research: Discontinuous criterion power analyses. *Annals of the International Communication Association*, 42, 75-92. doi: 10.1080/23808985.2018.1459198

Hopkins, D. J., & King, G. (2010). A method of automated nonparametric content analysis for social science. *American Journal of Political Science*, 54, 229–247. Doi:10.1111/j.1540-5907.2009.00428.x

- Hovy, E., & Lavid, J. (2010). Towards a “science” of corpus annotation: A new methodological challenge for corpus linguistics. *International Journal of Translation*, 22, 13-36.
- Krippendorff, K. (2004). Reliability in content analysis: Some common misconceptions and recommendations. *Human Communication Research*, 30, 411–433.
doi:10.1111/j.1468-2958.2004.tb00738.x
- Krippendorff, K. (2008). Validity. In W. Donsbach (Ed.), *The international encyclopedia of communication*. Hoboken, NJ: Blackwell Publishing.
doi:10.1002/9781405186407.wbiecv001
- Krippendorff, K. (2013). *Content analysis: An introduction to its methodology* (3rd ed.). Thousand Oaks, CA: Sage.
- Lease, M. (2011, August). On quality control and machine learning in crowdsourcing. In *Workshops at the Twenty-Fifth AAAI Conference on Artificial Intelligence*. Retrieved at: <https://www.aaai.org/ocs/index.php/WS/AAAIW11/paper/viewFile/3906/4255>
- Leemann, L., & Wasserfallen, F. (2017). Extending the use and prediction precision of subnational public opinion estimation. *American Journal of Political Science*, 61, 1003–1022. doi:10.1111/ajps.12319
- Lewis, S. C., Zamith, R., & Hermida, A. (2013). Content analysis in an era of big data: A hybrid approach to computational and manual methods. *Journal of Broadcasting & Electronic Media*, 57, 34–52. doi:10.1080/08838151.2012.761702
- Lind, F., Gruber, M., & Boomgaarden, H. G. (2017). Content analysis by the crowd: Assessing the usability of crowdsourcing for coding latent constructs. *Communication Methods and Measures*, 11, 191–209. doi:10.1080/19312458.2017.1317338
- Lowe, W., & Benoit, K. (2013). Validating estimates of latent traits from textual data using human judgment as a benchmark. *Political Analysis*, 21, 298–313.

IN VALIDATIONS WE TRUST

2

doi:10.1093/pan/mpt002

Muddiman, A., McGregor, S. C., & Stroud, N. J. (2018). (Re) claiming our expertise: Parsing

large text corpora with manually validated and organic dictionaries. *Political*

Communication, Online first, doi:10.1080/10584609.2018.1517843.

Rooduijn, M., & Pauwels, T. (2011). Measuring populism: Comparing two methods of

content analysis. *West European Politics*, 34, 1272–1283.

doi:10.1080/01402382.2011.616665

Scharkow, M. (2013). Thematic content analysis using supervised machine learning: An

empirical evaluation using German online news. *Quality & Quantity*, 47, 761–773.

doi:10.1007/s11135-011-9545-7

Scharkow, M., & Bachl, M. (2017). How measurement error in content analysis and self-

reported media use leads to minimal media effect findings in linkage analyses: A

simulation study. *Political Communication*, 34, 323–343.

doi:10.1080/10584609.2016.1235640

Young, L., & Soroka, S. (2012). Affective news: The automated coding of sentiment in

political texts. *Political Communication*, 29, 205–231.

doi:10.1080/10584609.2012.671234

IN VALIDATIONS WE TRUST

Table 1.

Results of the Systematic Literature Review, Study 1.

Method Used	<i>N</i>	Refer to Gold Standard	Report Reliability	Refer to Validation	Report Validation measures
Dictionary	55	29 (53%)	9 (16%)	27 (49%)	21 (38%)
Supervised ML	18	16 (89%)	10 (56%)	15 (83%)	12 (67%)
Total	73	45 (62%)	19 (26%)	42 (58%)	33 (45%)

Note: Percentages refer to share of articles using mentioned method.

Table 2.

Decision Scenarios Based on Observed vs. True Classification Performance, Study 2.

Observed performance	True classification performance	
	Below threshold	Above threshold
Below threshold	True Negative	False Negative (Type II)
Above threshold	False Positive (Type I)	True Positive

IN VALIDATIONS WE TRUST

3



Figure 1. Percentage of decision error and relative bias in F1-score (over 1000 simulations per each scenario), Naïve Bayes classifier. Results of GLM classifier and dictionary approaches are presented in online appendix.

Note: Upper panel = Proportion of cases (each simulation run) incorrectly conclude on classification performances. Lower panel = Relative bias in F1 scores among 1000 replications, with their median and 95% percentile confidence intervals.

Online Appendix

Online Appendix For:

In Validations We Trust? The Impact of Imperfect Human Annotations as a Gold Standard on the Quality of Validation of Automated Content Analysis

1. Variables coded in Study 1, detailed coding instructions, and reliability estimates

Following variables utilized in Study 1 were coded by 5 trained coders.

Variable	Definition & Coding instructions	Reliability
Relevance	Whether empirical text analysis is conducted and reported (Yes = 1, No = 0)	Alpha = 1
Method Used	1 = Search string based / Dictionary Approach 2 = Machine Learning 3 = Topic Modeling (excluded from further analysis) 4 = Other (excluded from further analysis)	Alpha = 1
Refer to gold standard	1 = Yes, a “gold standard” is used, and info is reported 0 = No is not used reported	Alpha = 1
Report reliability	Whether intercoder-reliability of human-coded materials are reported? (1 = Yes, reported, 0 = Not reported)	Alpha = 1
Refer to validation / Report validation measures	Whether validation of automated procedures are mentioned, and if so, whether either one of validation metrics (e.g., Recall, Sensitivity, Precision, Accuracy, F1, or other measures) is reported? (1 = Yes, mentioned, 0 = Not mentioned)	Alpha = .750

A total of five highly qualified coders tested the initial coding scheme by independently coding 10 sample articles (approximately 5% of the total retrieved sample) and collectively discussed any coding problems and disagreement. Coding instructions were iteratively revised until the coding schemes would produce reliable results.

Online Appendix

2. Detailed Setup of MC simulations

Supervised ML Scenario. We set three independent variables to be sampled from a multivariate normal distribution, with a randomly generated variance-covariance matrix Σ for each simulation run. This ensures that idiosyncratic values of the covariance matrix do not skew the overall results of the simulation. The true values of y (which is the binary variable) are then sampled from a Binomial distribution, with the probability parameter having a very simple linear functional form as follows:

$$y \sim \text{Bernoulli}(\pi)$$

$$\pi = \text{logistic}(\mu)$$

$$\mu = \mathbf{X}\beta + \epsilon$$

with ϵ being Gaussian noise added to ensure that each simulation run is not completely deterministic. The β , the true population parameter, was fixed throughout the simulation runs (specifically, $\beta_1 = 0.5$, $\beta_2 = 0.2$, and $\beta_3 = 0.6$, which were randomly chosen).

Dictionary-based Scenario. For a dictionary (i.e., bag-of-words) method, we assume a very similar approach as discussed above, but assume values of independent variables that were sampled from a Categorical distribution (i.e., a discrete value range from -5 to 5), where they represent some “features” of given textual data (e.g., a word or N-grams). Since this requires discrete rather than continuous values, we use a slightly different setup as follows:

$$\pi_{neg} \sim \text{Dirichlet}(N, \alpha_{neg})$$

$$\pi_{pos} \sim \text{Dirichlet}(N, \alpha_{pos})$$

$$\mathbf{X}_{k \in K}^T \sim \text{Categorical}(N, [\pi_{neg} \text{ or } \pi_{pos}])$$

with N being the total number of observations, and α_{neg} and α_{pos} being hyper-parameters governing the shapes of the categorical distribution, and X_k being a set of independent variables (with K being the number of textual features). The two negative and positive Dirichlet priors π_{neg} and π_{pos} were randomly selected for each independent variable,

Online Appendix

effectively treating such independent variables as systematic, recurring “features” of given textual data based on which y is generated in a similar fashion as in Equations (1). Yet for the dictionary-based approach, the vector β was extended to $K = 5$ and their β values were fixed to 0.2. This enables us to better approximate the multidimensionality of textual data, while treating y effectively as a function of the simple sum of the chosen textual features.

For Peer Review Only

Online Appendix




Figure A1. Overall classification accuracy against true value across MC simulation conditions, Study 2 (reference line is the overall mean).

Online Appendix



Figure A2. Percentage of decision error and relative bias in F1-score (over 1000 simulations per each scenario), GLM classifier.

Note: Upper panel = Proportion of cases (each simulation run) incorrectly conclude on classification performances. Lower panel = Relative bias in F1 scores among 1000 replications, with their median and 95% percentile confidence intervals.

Online Appendix



Figure A3. Percentage of decision error and relative bias in F1-score (over 1000 simulations per each scenario), Bag-of-words approach.

Note: Upper panel = Proportion of cases (each simulation run) incorrectly conclude on classification performances. Lower panel = Relative bias in F1 scores among 1000 replications, with their median and 95% percentile confidence intervals.